

 report_pdf.md

Training a Smart Cab

You can find all the relevant code for running this PyGame application and the smartcab navigating through it using Q-learning algorithm [here](#).

Installation

Python 2 Mac users

1. Install dependencies `brew install sdl sdl_ttf sdl_image sdl_mixer portmidi`
2. Install PyGame through Conda `conda install -c https://conda.anaconda.org/quasiben pygame`

Python 3 Mac userss

1. Create environment for Python 2.7 `conda create -n py27 python=2.x ipykernel`
2. Activate source `source activate py27`
3. Install dependencies `brew install sdl sdl_ttf sdl_image sdl_mixer portmidi`
4. Install PyGame through Conda `conda install -c https://conda.anaconda.org/quasiben pygame`

Implement a Basic Driving Agent

Code

There is only a change of code for `action` to choose an action randomly through the list `valid_actions = [None, 'forward', 'left', 'right']`. <https://gist.github.com/ritchieng/3a0a813e507f0e0ac68ebc9644fd71ac.js>

Trial 0: Reached Destination, Exceeded Time Limit

<https://gist.github.com/ritchieng/3debad4a7c75b64514ebb8d4c0790dc4.js>

Trial 1: Reached Destination, Within Time Limit

<https://gist.github.com/ritchieng/4f0948e859640ed4a38a08b47a83c7ed.js>

Observations

- If you can see on the first trial, Trial 0, the car took more than double the time of the given deadline of 25 moves.
- The car does eventually reach there with a few caveats:
 - i. It took a long time.
 - ii. It clashed with other cars.
 - iii. It made illegal moves (not obeying the traffic lights).
- Once in awhile, the car does reach the destination before the time is up.

Inform the Driving Agent

Appropriate States

- I have created tuples for the states as they are hashable.
- State
 - State means the environment that the smart cab encounters at every intersection.
 - This could be the:
 - Status of traffic lights (green or red)
 - Status of traffic at that intersection
 - Deadline
 - Next waypoint.
 - We have to decide what inputs you would like to include to define the state at the intersections.
- Relevant inputs:
 - Status of traffic lights: red and green.
 - This input is relevant as we need to obey the rules to reach our destination.
 - Next waypoint
 - This is relevant too because it concerns where the car should go.
- Irrelevant inputs:
 - Status of traffic at that intersection
 - This is a situation where there are few cars.
 - As such, the values for 'oncoming', 'left', and 'right' should almost always be 0.
 - We can leave this input out to simplify our learning process.
 - Deadline
 - I would normally include this.
 - However, there seem to be no rewards for reaching early. Only rewards for obeying traffic rules.
 - Hence, it is not worthwhile to include Deadline as an input as this would drastically increase the number of states we need to train on.

Number of States

- As such, we would have states that factor in next_waypoint (Left, Right and Forward) and light (red or green)
 - This would result in $3 \text{ (next_waypoint)} \times 2 \text{ (light)} = 6 \text{ states}$
- 6 states per position seem like a reasonable number given that the goal of Q-Learning is to learn and make informed decisions about each state.
 - This is because we have to understand that we face an exploration-exploitation dilemma here that is a fundamental trade-off in reinforcement learning.
 - This seems like a good balance of exploration and exploiting 6 states.

Code

<https://gist.github.com/ritchieng/905f12bf65265331f0e051541379c767.js>

Implement a Q-Learning Driving Agent

Estimating Q from Transitions: Q-learning Equation

- Personal notes:

Estimating Q from Transitions: Q-learning Equation

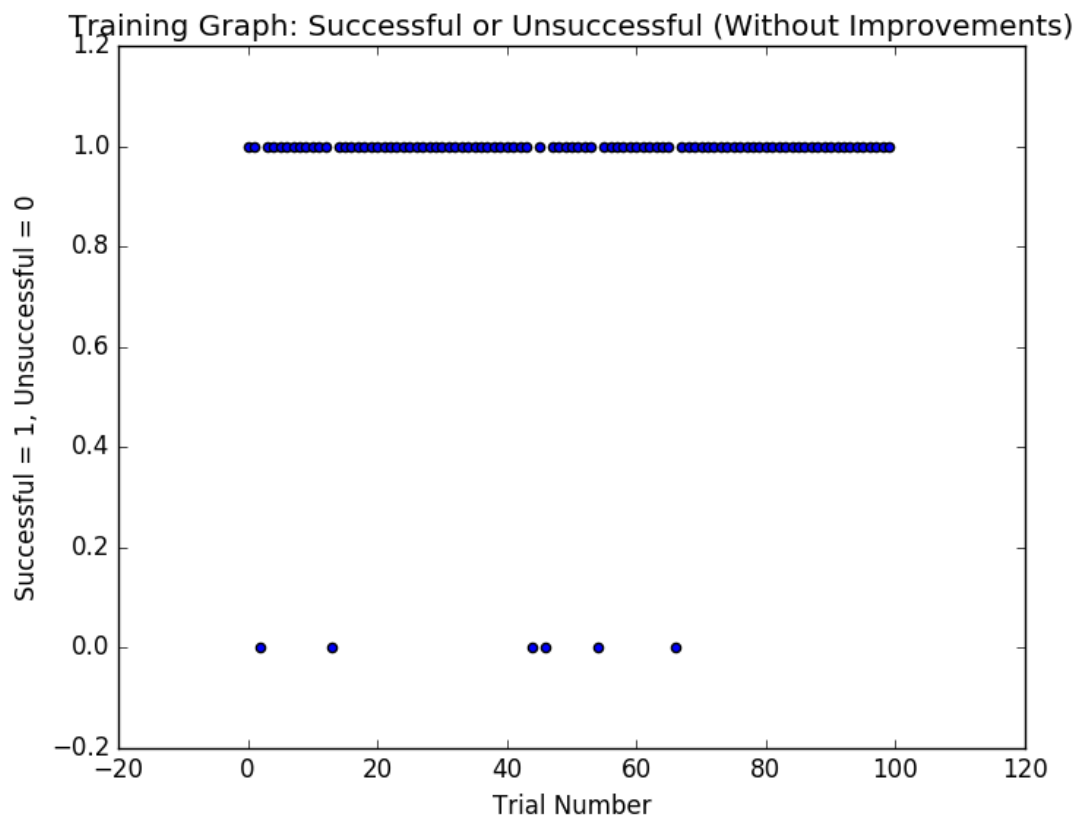
- $$\hat{Q}(s, a) \leftarrow_{\alpha_t} r + \gamma \max_{a'} \hat{Q}(s', a')$$
- **Intuition:** imagine if we've an estimate of the Q function.
 - We will update by taking the state and action and moving a bit.
 - We have the reward and discount the maximum of the utility of the next state.
- This represents the **utility function**.
 - $$r + \gamma \max_{a'} \hat{Q}(s', a')$$
- This represents the **utility of the next state**.
 - $$\max_{a'} \hat{Q}(s', a')$$
- α_t is the **learning rate**.
 - $V \leftarrow_{\alpha_t} X$
 - $V \leftarrow (1 - \alpha_t)V + \alpha X$
 - When $\alpha = 0$, you would have no learning at all where your new value is your original value, $V=V$.
 - When $\alpha = 1$, you would have absolute learning where you totally forget your previous value, V leaving your new value $V = X$.

Parameters Initiated

- Alpha (learning rate), is arbitrarily set at 0.5.
- Gamma (discount rate), is arbitrarily set at 0.5.
- Epsilon (randomness probability), is arbitrarily set such that it is 1%.
- Q initial values set at 1

Results

- The smart cab reaches the destination more frequently.
- Also, the smart cab takes fewer moves to reach to the destination.
 - The final trial took only 8 steps.
- Moreover, as you can see, we've achieved a success rate of 94% with a random assignment of parameters



Trial 100 results

<https://gist.github.com/ritchieng/02a2dd735ff4b13dcccfeb45fb4e07fe3.js>


Code

<https://gist.github.com/ritchieng/a43b7c188f083bb731efc2e78bd2ec4f.js>

Improve the Q-Learning Driving Agent

Parameter Tuning

- With trial and error, it seems the following parameters allow the agent to perform best.
 - Alpha (learning rate): 0.3
 - Gamma (discount rate): 0.3
 - Initial Q = 1
 - Success rate: 100%



Final trial results

<https://gist.github.com/ritchieng/c6f75bad8426f013310e5598a322391a.js>

Optimal Policy

- The agent does reach to the final absorbing states in the minimum possible time while incurring minimum penalties.
- Yet, it still does violate some traffic rules.
- The optimal policy would be:
 - Minimum possible time.
 - Obey all traffic rules.
 - No clashes with other cars.